

lintsampler: Easy random sampling via linear interpolation

Aneesh P. Naik¹ and Michael S. Petersen¹

¹ Institute for Astronomy, University of Edinburgh, UK ¶ Corresponding author

DOI: [10.xxxxxx/draft](https://doi.org/10.xxxxxx/draft)

Software

- [Review](#)
- [Repository](#)
- [Archive](#)

Editor: [Open Journals](#)

Reviewers:

- [@openjournals](#)

Submitted: 01 January 1970

Published: unpublished

License

Authors of papers retain copyright and release the work under a Creative Commons Attribution 4.0 International License ([CC BY 4.0](#)).

Summary

lintsampler provides a Python implementation of a technique we term ‘linear interpolant sampling’: an algorithm to efficiently draw pseudo-random samples from an arbitrary PDF. First, the PDF is evaluated on a grid-like structure. Then, it is assumed that the PDF can be approximated between grid vertices by the (multidimensional) linear interpolant. With this assumption, random samples can be efficiently drawn via inverse transform sampling (Devroye, 1986).

lintsampler is primarily written with numpy, drawing some additional functionality from scipy. Under the most basic usage of lintsampler, the user provides a PDF function and some parameters describing a grid-like structure to the LintSampler class, and is then able to draw samples via the sample method. Additionally, there is functionality for the user to set the random seed, employ quasi-Monte Carlo sampling, or sample within a premade grid (DensityGrid) or tree (DensityTree) structure.

Statement of need

Below is a (non-exhaustive) list of ‘use cases’, i.e., situations where a user might find lintsampler (and/or the the linear interpolant sampling algorithm underpinning it) to be preferable over other random sampling techniques.

In all of these use cases, it is assumed that the dimension of the problem is not too high. lintsampler works by evaluating a given PDF on the nodes of a grid (or grid-like structure, such as a tree), so the number of evaluations (and memory occupancy) grows exponentially with the number of dimensions. As a consequence, many of the efficiency arguments given for lintsampler below don’t apply to higher dimensional problems. We probably wouldn’t use lintsampler in more than 6 dimensions, but there is no hard limit here: the question of how many dimensions is too many will depend on the problem at hand.

A second assumption made below is that the target PDF the user wishes to sample from does not have its own exact sampling algorithm (such as the Box-Muller transform for a Gaussian PDF). The power of lintsampler lies in its applicability to arbitrary PDFs for which tailor-made sampling algorithms are not available.

Use Cases

1. Expensive PDF

If the PDF being sampled from is expensive to evaluate and a large number of samples is desired, then lintsampler might be the most cost-effective option. This is because lintsampler does not evaluate the PDF for each sample (as would be the case for importance sampling, rejection sampling or MCMC), but on the nodes of the user-chosen grid. Particularly in a low-dimensional

39 setting where the grid does not have too many nodes, this can mean far fewer PDF evaluations.
40 This point is demonstrated in the [first example notebook](#) in the `lintsampler` docs.

41 Similarly, there might be situations where the user is not so concerned about strict statistical
42 representativeness but wants to generate a huge number of samples from a target PDF with
43 the least possible computational cost (such as e.g., generating realistic point cloud distributions
44 in video game graphics). They can use `lintsampler` with a coarse grid (so minimal PDF
45 evaluations), then `sample()` to their heart's content.

46 2. Multimodal PDF

47 If the target PDF has a highly complex structure with multiple, well-separated modes, then
48 `lintsampler` might be the *easiest* option (in terms of user configuration). In such scenarios,
49 MCMC might struggle unless the walkers are carefully preconfigured to start near the modes.
50 Similarly, rejection sampling or importance sampling would be highly suboptimal unless the
51 proposal distribution is carefully chosen to match the structure of the target PDF. With
52 `lintsampler`, the user need only ensure that the resolution of their chosen grid is sufficient to
53 resolve the PDF structure, and so the setup remains straightforward. This is demonstrated in
54 the [second example notebook](#) in the `lintsampler` docs.

55 It is worth noting that in these kinds of complex, multimodal problems, a single fixed grid
56 might not be the most cost-effective sampling domain. For this reason, `lintsampler` also
57 provides simple functionality for sampling over multiple disconnected grids.

58 3. PDF with large dynamic range

59 If the target PDF has a very large dynamic range, then the `DensityTree` object provided
60 by `lintsampler` might be an effective solution. Here, the PDF is evaluated not on a fixed
61 grid, but on the leaves of a tree which is refined such that regions of concentrated probability
62 are more finely resolved. Such an example is shown in the [third example notebook](#) in the
63 `lintsampler` docs.

64 4. Noise needs to be minimised

65 In Quasi-Monte Carlo (QMC) sampling, one purposefully generates more 'balanced' (and thus
66 less random) draws from a target PDF, so that sampling noise decreases faster than $N^{-1/2}$.
67 `lintsampler` allows easy quasi-Monte Carlo sampling with arbitrary PDFs. We are not aware
68 of such capabilities with any other package. We give an example of using `lintsampler` for
69 QMC in the [fourth example notebook](#) in the `lintsampler` docs.

70 'Real World' Example

71 Any one of the four use cases above would serve by itself as a sufficient case for choosing
72 `lintsampler`, but here we give an example of a real-world scenario that combines all of the
73 use cases. It is drawn from our own primary research interests in computational astrophysics.

74 Much of computational astrophysics consists of large-scale high performance computational
75 simulations of gravitating systems. For example, simulations of planets evolving and interacting
76 within a solar system, simulations of stars interacting within a galaxy, or vast cosmological
77 simulations in which a whole universe is grown *in silico*.

78 There exists a myriad of codes used to run these simulations, each using different algorithms to
79 solve the governing equations. One class of simulation code that has gained much attention in
80 recent years is the class of code employing basis function expansions. In these codes, the matter
81 density at any point in space is represented by a sum over basis functions (not unlike a Fourier
82 series), with the coefficients in the sum changing over space and time. In these simulations,
83 the matter comprising the system is represented everywhere as a smooth, continuous fluid, but
84 for many applications and/or downstream analyses of the simulated system, one needs to

85 instead represent the system as a set of discrete particles. These particles can be obtained by
86 drawing samples from the continuous density field.

87 This is a scenario that satisfies all four of the use cases list above. To explain further: - The
88 PDF we are sampling from (i.e., the basis expansion representation of the matter density
89 field) can be expensive to evaluate if a large number of terms is included in the sum. - The
90 PDF can be highly multimodal when the system we are simulating comprises many distinct
91 gravitating substructures, such as stellar clusters. - The PDF can have a large dynamic range.
92 Astrophysical structures such as galaxies and dark matter 'haloes' often have power-law density
93 profiles, such as the Navarro-Frenk-White profile. Further complicating the issue is that a
94 typical dark matter halo will host several 'subhaloes', which in turn might host 'subsubhaloes',
95 and so on. In short, a range of spatial scales needs to be resolved. - If the particle set being
96 sampled is to be used for further simulation, it can be helpful to draw as 'noiseless' a sample
97 as possible for reasons of numerical stability.

98 For these reasons, this kind of astrophysical simulation code provides an excellent example
99 of a 'real world' application for `lintsampler`. Here, one would cover the simulation domain
100 with a `DensityTree` instance (or several instances, one for each primary structure), call the
101 `refine` method to better resolve the high-density regions, then feed the tree to a `LintSampler`
102 instance and call `sample` to generate particles. The `qmc` flag can be passed to the sampler in
103 order to employ Quasi-Monte Carlo sampling.

104 Features

105 Although `lintsampler` is written in pure Python, making the code highly readable, the
106 methods make extensive use of numpy functionality to provide rapid sampling. After the
107 structure spanning the domain has been constructed, sampling proceeds with computational
108 effort scaling linearly with number of sample points.

109 We provide two methods to define the domain, both optimised with numpy functionality for
110 efficient construction. The `DensityGrid` class takes highly flexible inputs for defining a grid.
111 In particular, the grid need not be evenly spaced (or even continuous) in any dimension; the
112 user can preferentially place grid elements near high-density regions. The `DensityTree` class
113 takes error tolerance parameters and constructs an adaptive structure to achieve the specified
114 tolerance. We also provide a base class (`DensityStructure`) such that the user could extend
115 the methods for spanning the domain.

116 Documentation for `lintsampler`, including example notebooks demonstrating a range of
117 problems, is available via a [readthedocs page](#). The documentation also has an extensive
118 explanation of the interfaces, including optimisation parameters for increasing the efficiency in
119 sampling.

120 Acknowledgements

121 We would like to thank Sergey Koposov for useful discussions. APN acknowledges funding
122 support from an Early Career Fellowship from the Leverhulme Trust. MSP acknowledges
123 funding support from a UKRI Stephen Hawking Fellowship.

124 References

125 Devroye, L. (1986). *Non-uniform random variate generation*. Springer-Verlag.